What is machine learning?

Artificial intelligence
Machine learning
What is machine learning?
How machine learning works
Types of machine learning
Supervised learning
Semi-supervised learning
Unsupervised learning
Reinforcement learning (RL)
Deep learning
Recurrent neural networks (RNNs)
Transformers
Mamba models
Machine learning use cases
Computer vision
Natural language processing (NLP)
Time series analysis
Image generation
Machine learning operations (MLOps)
Machine learning libraries

Author
Jim Holdsworth
Senior Technical Writer

What is machine learning?

Machine learning (ML) is a branch of artificial intelligence (AI) focused on enabling computers and machines to imitate the way humans do, that is, gradually improve their accuracy through exposure to more data. Machine learning algorithms are used in a variety of applications, such as email filtering and detection of network intruders, or software that is designed to analyze raw data and predict or classify what it is.

Machine learning begins with patterns that the AI model can recognize and use to make predictions. Data scientists can program certain rules or tasks into machine learning models to "guide" them in their learning process, but the model should be able to "learn" and make decisions based on the data it is exposed to.

Deep learning and neural networks are both forms of machine learning that mimic the way human brains operate. Deep learning models are made of many layers of neural networks that are trained on large quantities of data. The more data they are exposed to, the more they can learn and improve their accuracy. Deep learning is a subset of machine learning, and machine learning is a subset of artificial intelligence.

Generative AI models, such as OpenAI's GPT-4 and Google's Gemini, are examples of deep learning models. They are trained on very large quantities of data and can generate text, images, and other types of content based on the patterns they have learned from the data they were trained on. Generative AI models can also be fine-tuned to perform specific tasks, such as answering questions or translating languages, with relatively small amounts of additional data.

Machine learning is a powerful technology that can be used to automate a wide variety of tasks and to provide insights into data that would be difficult or impossible for humans to see on their own. However, it is important to be aware of the limitations of machine learning and to use it responsibly. For example, machine learning models can be biased if they are trained on biased data, and they can be used to make decisions that are unfair or discriminatory. It is important to be aware of these risks and to take steps to mitigate them when using machine learning.

Machine learning vs. artificial intelligence

Artificial intelligence (AI) is a broad term that refers to the ability of machines to perform tasks that normally require human intelligence, such as reasoning, learning, and problem-solving.

For example, imagine that a self-driving car encounters a pedestrian. Human drivers know that they must stop their car to avoid hitting the pedestrian. A self-driving car must be programmed with the ability to recognize pedestrians and to stop the car when it sees one. This is a task that requires intelligence, and it is one that machine learning can help with.

For example, machine learning can be used to train the self-driving car to recognize pedestrians by showing it many images of pedestrians. The more images of pedestrians the self-driving car sees, the better it will be at recognizing pedestrians and stopping the car when it sees one.

Machine learning is a powerful technology that can be used to automate a wide variety of tasks and to provide insights into data that would be difficult or impossible for humans to see on their own. However, it is important to be aware of the limitations of machine learning and to use it responsibly. For example, machine learning models can be biased if they are trained on biased data, and they can be used to make decisions that are unfair or discriminatory. It is important to be aware of these risks and to take steps to mitigate them when using machine learning.

Newsletter
Join over 100,000 subscribers who read the latest news in tech
Keep up with the latest advice, insights, tutorials, tips, tricks, and more in our weekly Think newsletter.
Enter your email address
Subscribe now
More than one newsletter? Manage your preferences.
Unsubscribe at any time.

How machine learning works

Machine learning models are fed large quantities of data to "learn" from and make predictions. For example, if a machine learning model is being trained to recognize pedestrians, it might be fed thousands of images of pedestrians and other objects. The more images of pedestrians the model sees, the better it will be at recognizing pedestrians and stopping the car when it sees one.

Machine learning models are typically trained using a process called supervised learning. In supervised learning, the model is given a set of data that has been labeled with the correct answers. For example, if a machine learning model is being trained to recognize pedestrians, the data might be labeled with the words "pedestrian" or "not pedestrian". The model then uses this labeled data to learn how to recognize pedestrians on its own.

Machine learning models also use parameters and optimization algorithms to learn from data. Parameters are variables that the model uses to make predictions. The model adjusts its parameters based on the data it is exposed to, in order to improve its accuracy. Optimization algorithms are used to find the best parameters for the model, given the data it is exposed to.

For example, if a machine learning model is being trained to recognize pedestrians, the parameters might be the size and shape of the pedestrians' heads and bodies. The model would then use these parameters to make predictions about whether or not it sees a pedestrian in an image. The optimization algorithm would then adjust the parameters based on the data the model is exposed to, in order to improve the model's accuracy in recognizing pedestrians.

For example, if a machine learning model is being trained to recognize pedestrians, the model might be fed thousands of images of pedestrians and other objects. The model would then use the parameters and optimization algorithm to learn how to recognize pedestrians on its own. The more images of pedestrians the model sees, the better it will be at recognizing pedestrians and stopping the car when it sees one.

Types of machine learning

There are several different types of machine learning models, each with its own strengths and weaknesses. Some of the most common types of machine learning models include:

- Supervised learning models are trained on a labeled dataset, meaning that the data has been labeled with the correct answers. This type of model is used for tasks such as classification and regression.
- Unsupervised learning models are trained on an unlabeled dataset, meaning that the data has not been labeled with the correct answers. This type of model is used for tasks such as clustering and dimensionality reduction.
- Reinforcement learning models are trained by being given a reward or penalty for each action they take. The model then learns how to take actions that will maximize the reward and minimize the penalty. This type of model is used for tasks such as game playing and robotics.
- Deep learning models are a type of neural network that is trained on large quantities of data. The more data the model is exposed to, the more it can learn and improve its accuracy. Deep learning models are used for tasks such as image recognition and natural language processing.
- Generative AI models are a type of deep learning model that is trained on large quantities of data and can generate text, images, and other types of content based on the patterns they have learned from the data they were trained on. Generative AI models can also be fine-tuned to perform specific tasks, such as answering questions or translating languages, with relatively small amounts of additional data.

Supervised learning

Supervised learning (also known as supervised machine learning) is defined by its use of labeled datasets to train algorithms that to classify data or predict outcomes accurately. As data is fed into the model, it adjusts its weights until it has been fitted appropriately, which occurs as part of the cross validation process. Supervised learning helps organizations solve a variety of real-world problems, such as classifying spam in a separate folder from your inbox. Some methods used in supervised learning include neural networks, naïve Bayes, linear regression, logistic regression, random forest, support vector machine (SVM), and more.

Supervised learning models can be very accurate if they are trained on high-quality data. The more data the model is exposed to, the more it can learn and improve its accuracy. However, supervised learning models can be biased if they are trained on biased data, and they can be used to make decisions that are unfair or discriminatory. It is important to be aware of these risks and to take steps to mitigate them when using supervised learning models.

Supervised learning models are used in a variety of applications, such as image recognition, natural language processing, and predictive analytics. Supervised learning models are also used in the development of artificial intelligence (AI) systems, such as self-driving cars and virtual assistants.

Semi-supervised learning

Semi-supervised learning offers a happy medium between supervised and unsupervised learning. During training, it uses a smaller labeled data set to guide classification and feature extraction from a larger, unlabeled data set. Semi-supervised learning can solve the problem of not having enough labeled data for a supervised learning algorithm to be fed after it's been trained on a smaller data set.

Unsupervised learning

Unsupervised machine learning uses algorithms to train on unlabeled data. The algorithm is then tasked with finding patterns in the data that it has been fed. Popular techniques in unsupervised learning include:

- Clustering algorithms are used to identify groups of similar data points in a dataset. For example, a clustering algorithm might be used to group customers by their purchase history or to group documents by their topic. Some popular clustering algorithms include k-means clustering, hierarchical clustering, and density-based spatial clustering of applications with noise (DBSCAN).
- Dimensionality reduction is a technique that is used to reduce the number of features in a dataset while preserving the most important information. This can be useful for improving the performance of machine learning models and for visualizing data. Some popular dimensionality reduction techniques include principal component analysis (PCA), t-distributed stochastic neighbor embedding (t-SNE), and uniform manifold approximation and projection (UMAP).
- Anomaly detection is a technique that is used to identify unusual or anomalous data points in a dataset. Anomaly detection can be used for fraud detection, intrusion detection, and quality control. Some popular anomaly detection techniques include isolation forest, one-class SVM, and local outlier factor (LOF).
- Association rule learning is a technique that is used to find relationships between variables in a dataset. For example, association rule learning might be used to find that customers who buy bread also tend to buy milk. This information can be used to make recommendations to customers or to optimize product placement in a store.

Unsupervised machine learning can be a powerful tool for discovering new knowledge and for making predictions about future events. However, it is important to be aware of the limitations of unsupervised machine learning and to use it responsibly. For example, unsupervised machine learning models can be biased if they are trained on biased data, and they can be used to make decisions that are unfair or discriminatory. It is important to be aware of these risks and to take steps to mitigate them when using unsupervised machine learning models.

Reinforcement learning (RL)

Reinforcement learning (RL) is a type of machine learning that is based on rewarding desired behaviors or punishing undesired behaviors. In this type of machine learning, an agent is trained to perform a task in an environment by receiving rewards for actions that move it closer to its goal and penalties for actions that move it farther from its goal. The agent learns from its experiences and adjusts its behavior to maximize the rewards it receives and minimize the penalties it receives.

Reinforcement learning is used in a variety of applications, such as robotics, game playing, and autonomous driving. For example, reinforcement learning might be used to train a robot to navigate a maze or to play a game of chess. The robot would be rewarded for actions that move it closer to the goal of navigating the maze or winning the game of chess, and penalized for actions that move it farther from the goal.

Reinforcement learning is a powerful technique that can be used to train agents to perform complex tasks in environments that are difficult for humans to navigate. However, it is important to be aware of the limitations of reinforcement learning and to use it responsibly. For example, reinforcement learning models can be biased if they are trained on biased data, and they can be used to make decisions that are unfair or discriminatory. It is important to be aware of these risks and to take steps to mitigate them when using reinforcement learning models.

Deep learning

Deep learning models are trained by using large sets of labeled data and neural network architecture learning features directly from the data without the need for human feature engineering. Deep learning models are able to "learn" from unstructured and unlabeled data by identifying which features in large datasets can be used to accurately cluster or classify data points. For example, deep learning can be used to classify a set of photos by their subject matter, such as cats, dogs, or landscapes, without the need for human labelers to first label each photo with the appropriate category.

Deep learning models are trained by using a large set of labeled data and neural network architectures that learn features directly from the data without the need for human feature extraction. Deep learning models are able to achieve state-of-the-art accuracy on many tasks, such as image recognition, natural language processing, and speech recognition. However, deep learning models can be computationally expensive to train and require large amounts of data to be effective.

Deep learning models are used in a variety of applications, such as self-driving cars, medical diagnosis, and fraud detection. For example, deep learning might be used to train a self-driving car to recognize pedestrians and other objects on the road, or to diagnose a medical condition from a patient's medical imaging. In the case of fraud detection, deep learning might be used to identify patterns in financial transactions that are indicative of fraud.

Deep learning models are a powerful technique that can be used to achieve state-of-the-art accuracy on many tasks. However, it is important to be aware of the limitations of deep learning and to use it responsibly. For example, deep learning models can be biased if they are trained on biased data, and they can be used to make decisions that are unfair or discriminatory. It is important to be aware of these risks and to take steps to mitigate them when using deep learning models.

Recurrent neural networks (RNNs)

Recurrent neural networks (RNN) are a type of neural network that is designed to process sequential data, such as text, speech, or time series data. RNNs are able to "remember" the previous inputs they have received and use this information to make predictions about the current input. This makes RNNs useful for tasks such as language modeling, machine translation, and speech recognition.

RNNs are trained using a process called backpropagation through time. This process involves feeding the network a sequence of data and then adjusting the weights of the network based on the error between the predicted output and the actual output. The weights are then updated to minimize the error in the future.

RNNs are used in a variety of applications, such as language modeling, machine translation, and speech recognition. For example, RNNs might be used to predict the next word in a sentence, translate a sentence from one language to another, or transcribe spoken words into text.

RNNs are a powerful technique that can be used to achieve state-of-the-art accuracy on many tasks. However, it is important to be aware of the limitations of RNNs and to use them responsibly. For example, RNNs can be biased if they are trained on biased data, and they can be used to make decisions that are unfair or discriminatory. It is important to be aware of these risks and to take steps to mitigate them when using RNNs.

Transformers

Transformers (or transformer models) are a type of neural network that has become very popular in the field of natural language processing (NLP). They are designed to process sequential data, such as text, and have been used to achieve state-of-the-art results on many NLP tasks, such as machine translation, text summarization, and question answering. Transformers are also used in other applications, such as speech recognition and image recognition.

Mamba models

Mamba is a newer neural network architecture that combines the strengths of two other key types of neural networks: state space models (SSMs) and transformers. SSMs are good at capturing the overall context of a sequence of data, while transformers are good at paying attention to the most relevant parts of the input data. Mamba is a hybrid model that uses both SSMs and transformers to process sequential data, such as text or time series data. The Mamba architecture is designed to be more efficient than traditional transformers, and it has been shown to achieve state-of-the-art performance on a variety of tasks, such as language modeling and machine translation.

Machine learning use cases

Machine learning models are used in a variety of ways, with new use cases being discovered all the time. Some of the most common use cases include:

Computer vision

Computer vision is a field of artificial intelligence that enables computers to see and interpret visual information, such as images and videos. Machine learning models are used in computer vision to perform tasks such as image classification, object detection, and facial recognition.

Natural language processing (NLP)

Natural language processing (NLP) is a field of artificial intelligence that enables computers to understand and process human language. Machine learning models are used in NLP to perform tasks such as sentiment analysis, machine translation, and text summarization.

Time series analysis

Time series analysis is a field of statistics that deals with analyzing data that changes over time. Machine learning models are used in time series analysis to perform tasks such as forecasting, anomaly detection, and trend analysis.

Image generation

Image generation is a field of artificial intelligence that enables computers to create new images that look like they were created by a human. Machine learning models are used in image generation to perform tasks such as image synthesis and image editing.

Machine learning operations (MLOps)

Developing machine learning models can be challenging. Operationalizing machine learning models can be even more challenging.

Machine learning operations (MLOps) is a set of practices that aims to deploy and maintain machine learning models in production reliably and efficiently. It includes activities such as data preparation, model training, model deployment, and monitoring. The goal of MLOps is to ensure that machine learning models are deployed and maintained in a way that is scalable, reliable, and secure.

Machine learning libraries

Machine learning libraries are collections of pre-written code that can be used to implement machine learning algorithms. These libraries provide a wide variety of functions and tools that can be used to perform tasks such as data preprocessing, model training, and model evaluation. Some popular machine learning libraries include TensorFlow, PyTorch, and scikit-learn.

TensorFlow is an open-source library developed by Google that is used for machine learning and artificial intelligence. It provides a wide variety of functions and tools that can be used to perform tasks such as data preprocessing, model training, and model evaluation. TensorFlow is known for its flexibility and scalability, and it is used by many companies and organizations for a variety of machine learning applications.

PyTorch is an open-source library developed by Facebook that is used for machine learning and artificial intelligence. It provides a wide variety of functions and tools that can be used to perform tasks such as data preprocessing, model training, and model evaluation. PyTorch is known for its ease of use and its support for dynamic computation graphs, which make it easier to debug and optimize machine learning models.

Scikit-learn is an open-source library developed by the scikit-learn community that is used for machine learning and artificial intelligence. It provides a wide variety of functions and tools that can be used to perform tasks such as data preprocessing, model training, and model evaluation. Scikit-learn is known for its simplicity and its support for a wide variety of machine learning algorithms.

There are many other machine learning libraries available, and the choice of which library to use will depend on the specific requirements of the project. It is important to be aware of the strengths and weaknesses of each library and to choose the one that is most suitable for the task at hand.

Subscribe to our weekly AI newsletter
Get the latest updates from industry experts, research breakthroughs, and more in our weekly AI newsletter.
Subscribe now
More than one newsletter? Manage your preferences.
Unsubscribe at any time.

Resources

Guide
The 2025 IDC guide to modernizing mainframe application development
Learn how to accelerate mainframe application development with AI-powered tools, boost productivity, and maintain high-quality code with confidence.
Read the report
Guide
Learning path: RAG essentials
Learn how to use retrieval-augmented generation (RAG) to connect large language models to your data for more accurate, relevant responses.
Explore the course
Guide
Learning path: Generative AI fundamentals on watsonx
Learn how to develop foundational knowledge of generative AI, large language models (LLMs), and prompt engineering.
Explore the course
Guide
How to scale, driving ROI with gen AI
Learn how to scale generative AI with a strategy that drives business value across your organization.
Read the guide

Related solutions

AI for business leaders
Learn how to leverage AI and machine learning to gain a competitive edge in your industry and drive business outcomes.
Explore AI for business leaders
AI for developers
Learn how to build, deploy, and manage AI models and applications across any cloud or environment.
Explore AI for developers
AI for IBM z and LinuxONE
Learn how to unleash the power of AI on your IBM z and LinuxONE platform to accelerate business insights and enhance security and compliance.
Explore AI for IBM z and LinuxONE

Take the next step

Watsonx.ai is the next generation enterprise studio for AI builders. Train, validate, tune and deploy AI models.
Start for free
Talk to an expert

Contributors
Author: Jim Holdsworth, Senior Technical Writer
Editors: Alexandra Jonker, Senior Editor; Ian Smalley, IBM Cloud Learn Hub Content Lead
Reviewers: Ken Hess, Senior Technical Writer; Ian Smalley, IBM Cloud Learn Hub Content Lead
Illustrations: Jacquelyn Cuccia, Senior Interactive Designer; Kim Whitaker, Senior Interactive Designer
Guest: Jeff Crume, Distinguished Engineer, Master Inventor, IBM Security; Martin Keen, Master Inventor, Technical Content Strategist
Last updated: 16 August 2024